

Homework 0

Alohomora

Shounak Naik
Masters in Robotics
WPI

I. INTRODUCTION

This report contains my understanding and implementation of the homework 0. The first phase is related to edge detection and the second phase is related to modelling neural networks.

II. PHASE 1: SHAKE MY BOUNDARY

In this section I will explain how I approached and solved the Pb-Lite related assignment. There are 4 filter banks generated before the actual implementation of the Pb-Lite algorithm.

A. DoG Filter Bank

For creating the Derivative of Gaussian, I took a 7x7 Gaussian matrix. I took a 3x3 dimensional Sobel. Convoluting Gaussian and Sobel results in an approximation of the Derivative of Gaussian. Once I had the DoG matrix, I used the OpenCV to rotate the matrix so that I could get different orientations of the same matrix. The Sobel Kernel used in my program is as follows:

$$Sobel = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

The Gaussian matrix used for the DoG was a ready-made 7x7 Gaussian matrix. For all other banks I have the following Gaussian equation 1 to create Gaussian matrices.

$$G(x, y) = \frac{1}{2\pi\sigma_x\sigma_y} \exp\left(-\frac{x^2}{2\sigma_x^2}\right) \exp\left(-\frac{y^2}{2\sigma_y^2}\right) \quad (1)$$

In Fig 1, you can see the DoG filter library obtained by me.

B. LM Filter Bank

For the LMS, the scales taken were 1, $\sqrt{2}$ and 2. For LML, the scales taken were $\sqrt{2}$, 2 and $2\sqrt{2}$. To generate the different orientations, I used OpenCV rotate function.

1) First Order Derivative:

The Gaussian used in this case has 2 different σ 's - σ_x and σ_y . Thus Equation 1 was used to create these special Gaussian matrices. Here $\sigma_x = 3\sigma_y$. After getting these Gaussian matrices, they were convolved with the Sobel filter to get the first order derivative Gaussian. Of course this is an approximation. The left side of the first 3 rows in Figure 2 and Figure 3 are the First Order matrices

2) Second Order Derivative:

Convolution between the Sobel Filter and the First Order




Fig. 1. Derivative of Gaussian Bank




Fig. 2. LM Small




Fig. 3. LM Large



Fig. 4. Gabor filter bank

Derivative matrix gives us the Second Order matrix. The right side of the first 3 rows in Figure 2 and Figure 3 are the Second Order matrices.

3) Laplacian of Gaussian:

LoG was obtained directly with the equation 2. LoG occurs at σ and 3σ . LoG can be seen in the first 8 images in the last row in Figure 2 and Figure 3.

$$L(x, y) = -\frac{1}{\pi\sigma^4} \left[1 - \frac{x^2 + y^2}{2\sigma^2} \right] \exp\left(-\frac{x^2 + y^2}{2\sigma^2}\right) \quad (2)$$

4) Gaussian:

Gaussian is obtained by 4 basic scales. These basic scales change according to the type of LM bank that we want to create (Small or Large). Gaussian maps can be seen in the last 4 images in the last row in Figure 2 and Figure 3.

C. Gabor Filters

These filters caught my eye because they also appear in the human visual system. These filters are created by modulating a sinusoidal wave onto the Gaussian kernel. There are 5 parameters - σ , γ , θ , λ and ψ . I experimented with the way sinusoidal wave can be modulated onto the Gaussian kernel. I tried adding *cosine* and *sin* components, but best results were obtained with using just the real component that is the *cosine* component. Further I only varied parameters of σ , θ and γ . I observed the best visual results are when the σ parameter is near 2. The Gabor filters created can be seen in Figure 4

D. Different Maps

Texton Map was generated by convolving grayscale images with all filter in all of the 3 filter banks mentioned above. I got a 3D tensor by doing the above operation. To quantize this tensor into 2D, I used the Kmeans clustering algorithm with number of cluster as 64. To do this flattened out the multiple 2D images in the 3D tensor. I passed this flattened out images as my input to the Kmeans algorithm. Once I got the Kmeans labels, I resized the flattened array into a 2D array. The resultant Texton map can be seen in 5.

A similar algorithm was used to create the Brightness and the Color Maps. Brightness map only had one channel and the Color Map had 3 channels. For both the Brightness and the Color Map, the number of clusters used was 16.




Fig. 5. Texton Map




Fig. 6. Half Disk Masks

E. Gradient Maps

1) *Half Disk Masks*: These masks are useful for calculating the Chi-square distance later. The way I created these half disc masks was by creating masks and then applying them on a fully black image. Firstly, I created a circular mask and then a rectangular mask of side equalling the diameter of the circle. Then I AND'd these 2 masks and then applied this mask to the black image. The Half Disk bank created can be seen in Figure 6

2) *Chi-Square Gradient*: This was the most tough and the most interesting part of Phase 1 for me. I understood the binning concept only when I understood what a histogram of an image is. I took the bin size as the number of clusters used to create the maps earlier. Firstly for every bin, I masked the map to only contain the bin number. Then I convolved pairs of opposing filters onto the masked map and then calculated the Chi-Square distance according to the Chi-square distance formula given in the assignment.

Since the tensor created from this algorithm was of 3 dimensions, I took a mean across the 3rd dimension - across the same pixel. This gave me a 2D output.

Since many values were 0's, the division led to some NaN values. I handled these cases by simply putting a 0 in their place. This was the most logical solution I could think of as putting 0 would just give us a black pixel and won't give us a false positive pixel. The warning does pop up when executed




Fig. 7. Texton Gradient




Fig. 8. Brightness Gradient

but that can be ignored as I handle the warning later.

Texton Gradient, Brightness Gradient and Color Gradient were created by using the same above concept. Texton Gradient can be seen in Figure 7. Brightness Gradient can be seen in Figure 8. Color Gradient can be seen in Figure 9

3) *Pb-Lite Final Output*: A mean over the Texton Gradient, Brightness Gradient and Color Gradient is taken. Canny and Sobel baseline is linearly weighted together with w_1 and w_2 being 0.5 both. w_1 and w_2 are weights for Canny and Sobel respectively. The final Pb-Lite output can be seen in Figure




Fig. 9. Color Gradient




Fig. 10. Pb-Lite Output

10. The figure is light and the edges can be seen when seen closely.

I believe using this algorithm lets us introduce texture, brightness and color components into the edge detection problem. This also reduces other irrelevant edges from the scene.

III. DEEP DIVE ON DEEP LEARNING

This section will be a walk through the Phase2 of the assignment. It will discuss the Convolutional Neural Network architecture that I built to tackle the multiclass classification of the CIFAR-10 dataset. I will also be discussing the theoretically improvements that can be made to the first Neural Network I build in this assignment.

A. Train your First Neural Network

All images in the CIFAR-10 dataset are $32 \times 32 \times 3$ and we are supposed to classify them into 10 classes. I went with a relatively simplistic CNN model to solve this problem. The architecture that I built is seen in Figure 11. I experimented with the number of Linear Layers - addition of another Linear layer actually dropped the accuracy significantly. (From 58 % in the current architecture to 50 % in the deeper network). I have trained my network for 10 epochs.

I have used the *torchsummary* package to summarize model architecture and parameters. My network has the following configurations:

- 1) Number of Parameters: 50,862 (all of them trainable)
- 2) Optimizer: Adam with $LR = 0.001$
- 3) Batch Size: 200

The Training Loss graph can be seen in 12. The Training Accuracy graph can be seen in 13. The Test Accuracy graph can be seen in 14. To generate the test accuracy plot, I loaded the model checkpoints at every epoch and calculated the test accuracy. I feel that instead of a test accuracy graph, a validation accuracy graph would be ideal. 3 splits (train, test and validation) would have been better instead of 2 splits (train and test).




Fig. 11. Primitive CNN




Fig. 12. Training Loss over Epochs

The test confusion matrix can be seen in Figure 15. The train confusion matrix can be seen in Figure 16. The test accuracy is 58.07 % and the train accuracy is 60.61 %.

B. Improving Accuracy of your Neural Network

I have tried to improve my primitive CNN I built earlier. To do this, I used the *torchvision transform* API to make a transformation composition. In this composition, I have first randomly rotated photos and then applied a random inversion of colours. This was done so that my network does not overfit on a specific colour too much. Finally, I normalized the images just before I fed them to my network. I also made some changes to my network - I added a Batch Normalization 2D to my Convolutional Network. I applied the BatchNorm just before the MaxPooling layers of my 2 convolutional layers. (See Figure 11)




Fig. 13. Training Accuracy over Epochs




Fig. 14. Test Accuracy over Epochs

```

Test Confusion Matrix
0% | 0/10000 [00:00<?, ?it/s]
[ 704 27 77 17 15 19 10 11 77 43] (0)
[ 46 640 22 15 6 17 10 7 44 193] (1)
[ 84 7 504 74 62 133 71 33 12 20] (2)
[ 33 12 110 353 30 317 72 32 10 31] (3)
[ 36 6 215 66 356 124 73 100 14 10] (4)
[ 11 5 107 128 23 635 22 48 6 15] (5)
[ 9 7 80 88 35 59 679 18 2 23] (6)
[ 26 0 69 46 48 145 7 620 4 35] (7)
[172 41 38 18 11 27 8 7 622 56] (8)
[ 64 91 30 19 5 21 12 20 44 694] (9)
(0) (1) (2) (3) (4) (5) (6) (7) (8) (9)
Accuracy: 58.07 %
  
```

Fig. 15. Test Confusion Matrix

```

Train Confusion Matrix
0% | 0/50000 [00:00<?, ?it/s]
[3562 114 396 89 74 79 32 59 373 222] (0)
[ 291 3355 94 56 38 72 64 46 222 762] (1)
[ 348 50 2716 362 319 577 307 174 70 77] (2)
[ 121 26 599 1900 134 1520 353 142 59 146] (3)
[ 233 27 1036 319 1945 555 310 458 47 70] (4)
[ 60 21 461 715 141 3196 107 208 31 60] (5)
[ 35 47 471 410 169 276 3398 55 32 107] (6)
[ 97 12 308 202 210 674 39 3293 16 149] (7)
[ 745 162 172 86 49 87 38 15 3363 283] (8)
[ 287 474 104 91 23 104 81 94 165 3577] (9)
(0) (1) (2) (3) (4) (5) (6) (7) (8) (9)
Accuracy: 60.61 %
  
```

Fig. 16. Train Confusion Matrix




Fig. 17. Training Loss over Epochs for Improved Network




Fig. 19. Test Accuracy over Epochs for Improved Network




Fig. 18. Training Accuracy over Epochs for Improved Network

```

Test Confusion Matrix
0%|          | 0/10000 [00:00<?, ?it/s]
[526 30 93 81 29 16 9 31 154 31] (0)
[ 48 451 34 61 9 44 27 41 128 157] (1)
[140 25 261 72 197 94 51 60 86 14] (2)
[ 37 30 86 250 84 228 69 76 84 56] (3)
[ 69 20 81 47 477 44 81 117 37 27] (4)
[ 43 29 50 140 70 456 29 77 71 35] (5)
[ 20 33 36 59 97 58 572 40 45 40] (6)
[ 64 40 42 71 101 103 22 436 43 78] (7)
[164 68 43 70 11 24 14 23 507 76] (8)
[ 44 168 12 109 13 38 27 55 95 439] (9)
(0) (1) (2) (3) (4) (5) (6) (7) (8) (9)
Accuracy: 43.75 %

```

Fig. 20. Test Confusion Matrix for Improved Network

IV. CONCLUSION

I was not at all familiar with the part of Computer Vision pertaining to Phase 1. I enjoyed solving the Phase 1 part as it was new and challenging. The first phase took me a lot of time which left me little time to complete the last part of Phase 2. I look forward to the upcoming assignments in the course.

ACKNOWLEDGMENT

I would like to thank the Course team of Computer Vision to allow me learn hands on various things about filters.

Since I have augmented the data using transformations, I increased the epochs to 30. Due to the addition of BatchNorm the parameters of my network increased to 50,916. Adam Optimizer already adjusts learning rate during the training process and thus I did not change my optimizer.

The Training Loss graph can be seen in 17. The Training Accuracy graph can be seen in 18. The Test Accuracy graph can be seen in 19. The test confusion matrix can be seen in Figure 20. The train confusion matrix can be seen in Figure 21. The test accuracy is 43.75 % and the train accuracy is 44.18 %.

C. Comparison

The augmentation steps I added to the training process negatively affected the accuracy of the classification task. I believe the augmentations might be too confusing for the model to understand and maybe simplifying the augmentation steps can benefit the accuracy metric. I also believe a higher learning rate might improve the accuracy as Adam and BatchNorm combined together would allow for a higher learning rate.

I did not implement the last part of the assignment due to lack of time. I would be implementing the Network architecture as soon as I get the time.

```

Train Confusion Matrix
0%|          | 0/50000 [00:00<?, ?it/s]
[2654 209 441 328 137 121 37 167 740 166] (0)
[ 264 2319 151 311 48 235 123 210 604 735] (1)
[ 624 107 1516 325 956 395 256 304 403 114] (2)
[ 234 164 375 1210 407 1175 420 330 421 264] (3)
[ 271 73 440 257 2486 230 308 551 201 183] (4)
[ 216 169 338 665 375 2126 186 396 352 177] (5)
[ 87 173 220 302 478 249 2792 219 270 210] (6)
[ 281 171 255 321 538 500 77 2240 197 420] (7)
[ 770 393 239 447 82 135 69 107 2439 319] (8)
[ 203 744 86 401 57 259 113 346 480 2311] (9)
(0) (1) (2) (3) (4) (5) (6) (7) (8) (9)
Accuracy: 44.186 %

```

Fig. 21. Train Confusion Matrix for Improved Network